

Kubernetes rolling updates, rollbacks and multi-environments

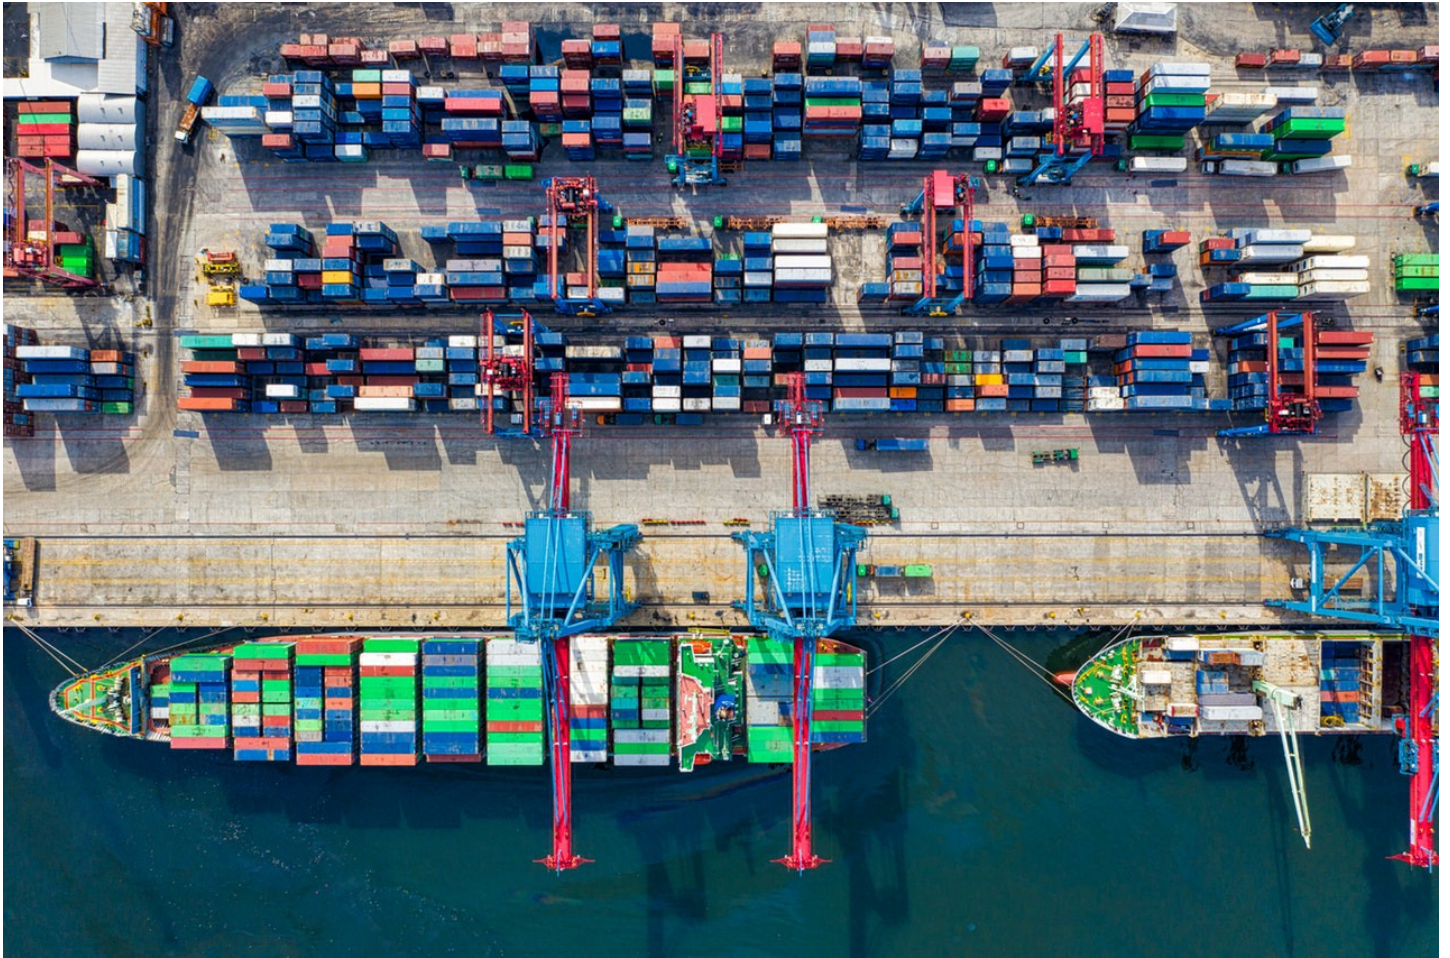


Photo by Tom Fisk from Pexels

On previous post (<https://itnext.io/deploy-an-app-on-kubernetes-gke-with-kong-ingress-letsencrypt-and-cloudflare-94913e127c2b>) we learned how to deploy an application with 2 micro-services (frontend and backend) to Kubernetes, using Kong Ingress, LetsEncrypt to provide TLS certificates and Cloudflare for proxy and extra security.

In this post we want to do some updates to our deployed application, roll them back in the case of errors and last but not least use multiple environments so we can test our application before deploying to production.

First I will re-deploy my original application. (I always delete un-used applications, no need to spend money on hosting them)

One nifty feature of `kubectl` is you can concatenate all resource files and apply them on bulk. So this is my file:

```
1  apiVersion: v1
2  kind: Namespace
3  metadata:
4    name: outsrc
5  ---
6  apiVersion: extensions/v1beta1
7  kind: Ingress
8  metadata:
9    name: outsrc-dev-ingress
10   namespace: outsrc
11   annotations:
12     kubernetes.io/ingress.class: kong
13     kubernetes.io/tls-acme: 'true'
14     cert-manager.io/cluster-issuer: letsencrypt-production
15 spec:
16   tls:
17     - secretName: outsrc-dev-tls
18     hosts:
19       - outsrc.dev
20   rules:
21     - host: outsrc.dev
22       http:
23         paths:
24           - path: /api
25             backend:
26               serviceName: service-backend
27               servicePort: 3000
28           - path: /
29             backend:
30               serviceName: service-frontend
31               servicePort: 3000
32   ---
33  apiVersion: v1
34  kind: Service
35  metadata:
36    name: service-frontend
37    namespace: outsrc
38    labels:
39      service: front
40  spec:
```

```
41     selector:
42       service: front
43   ports:
44     - port: 3000
45       protocol: TCP
46       targetPort: 3000
47   ---
48   apiVersion: apps/v1
49   kind: Deployment
50   metadata:
51     name: deployment-frontend
52     namespace: outsrc
53     labels:
54       service: front
55   spec:
56     replicas: 2
57     selector:
58       matchLabels:
59         service: front
60     template:
61       metadata:
62         labels:
63           service: front
64       spec:
65         containers:
66           - name: frontend-container
67             image: 'gcr.io/outsrc/outsrc-demo-front:1.0.0'
68             imagePullPolicy: Always
69             ports:
70               - containerPort: 3000
71             env:
72               - name: API_URL
73                 value: 'https://outsrc.dev/api'
74               - name: PORT
75                 value: '3000'
76   ---
77   apiVersion: v1
78   kind: Service
79   metadata:
80     name: service-backend
81     namespace: outsrc
82     labels:
83       service: back
84   spec:
85     selector:
86       service: back
87     ports:
88       - port: 3000
```

```

88     - port: 3000
89     protocol: TCP
90     targetPort: 3000
91 ---
92 apiVersion: apps/v1
93 kind: Deployment
94 metadata:
95   name: deployment-backend
96   namespace: outsrc
97   labels:
98     service: back
99 spec:
100   replicas: 2
101   selector:
102     matchLabels:
103       service: back
104   template:
105     metadata:
106       labels:
107         service: back
108     spec:
109       containers:
110         - name: backend-container
111           image: 'gcr.io/outsrc/outsrc-demo-back:1.0.0'
112           imagePullPolicy: Always
113           ports:
114             - containerPort: 3000
115           env:
116             - name: PORT
117               value: '3000'

```

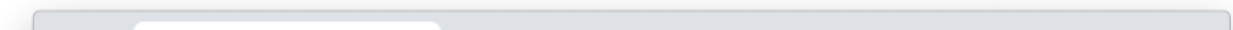
to apply it.

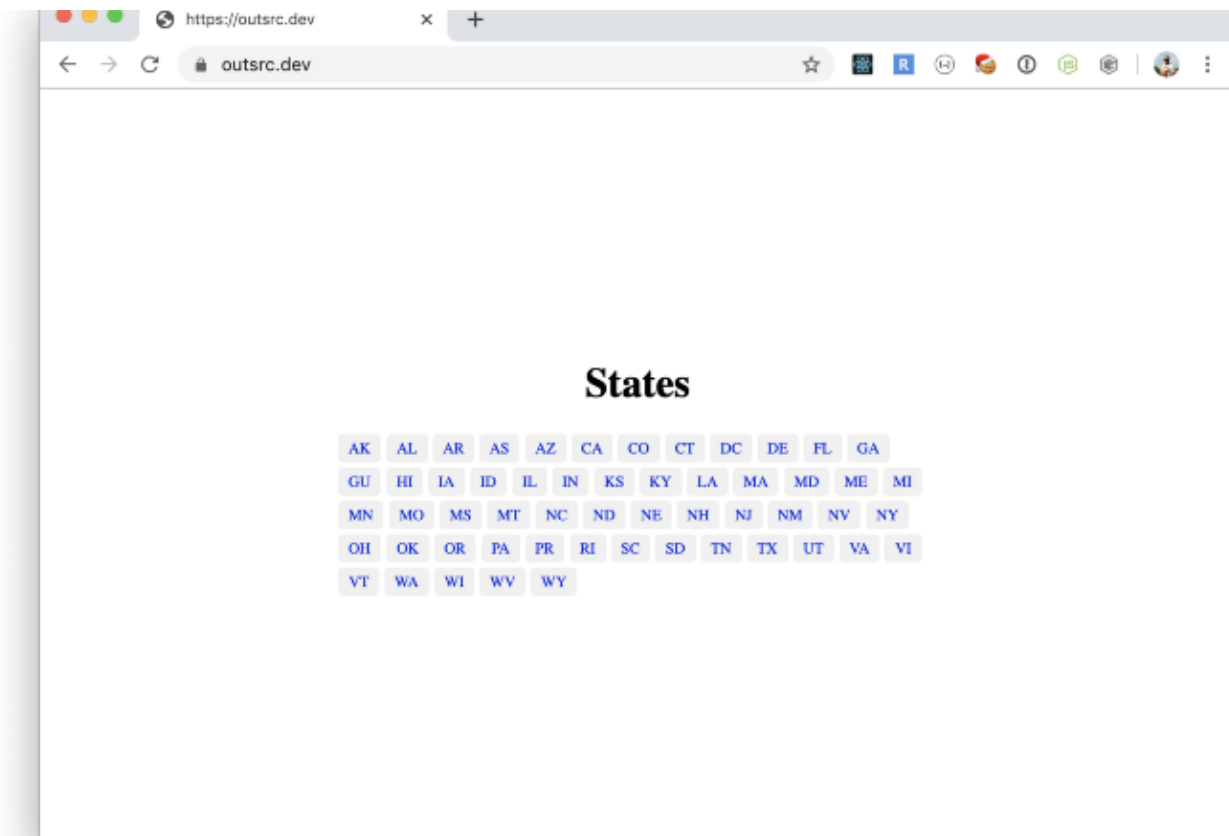
```

$ kubectl apply -f outsrc.yml
namespace/outsrc created
ingress.extensions/outsrc-dev-ingress created
service/service-frontend created
deployment.apps/deployment-frontend created
service/service-backend created
deployment.apps/deployment-backend created

```

And after a couple seconds (container running, TLS certs emitted) we have the application back online:





outsrc.dev deployed to Kubernetes

So far all good. Now, we need to make some changes to the Application. We will add a map page containing the US map, we will link the map from both the main page and the state page. After adding the code for this feature and dockerize it:

```
$ docker build -t outsrc-demo-front .  
...  
  
$ docker tag outsrc-demo-front:latest gcr.io/outsrc/outsrc-demo-front:1.1.0  
  
$ docker push gcr.io/outsrc/outsrc-demo-front:1.1.0  
...
```

Notice the version is different, We will use docker image tags to do a rolling update.

Rolling updates

Every time we want to update the application we deployed on our kubernetes cluster we change our deployment resource files and update them. Each time a change is detected a rolling update will be performed.

To avoid downtime kubernetes will update each replica of our running container one by one and re-routing the services on top.

To check updates history:

```
$ kubectl rollout history deployment/deployment-frontend
deployment.extensions/deployment-frontend
REVISION  CHANGE-CAUSE
1          <none>
```

Lets update the frontend container image tag:

```
$ kubectl set image deployment/deployment-frontend frontend-
container=gcr.io/outsrc/outsrc-demo-front:1.1.0
deployment.extensions/deployment-frontend image updated
```

We could also modify the deployment resource file, change the image tag and apply it via `kubectl` . This is my preferred way to handle updates, since it keeps the source of truth on the resource descriptor files.

Watch the update:

```
$ kubectl rollout status -w deployment/deployment-frontend

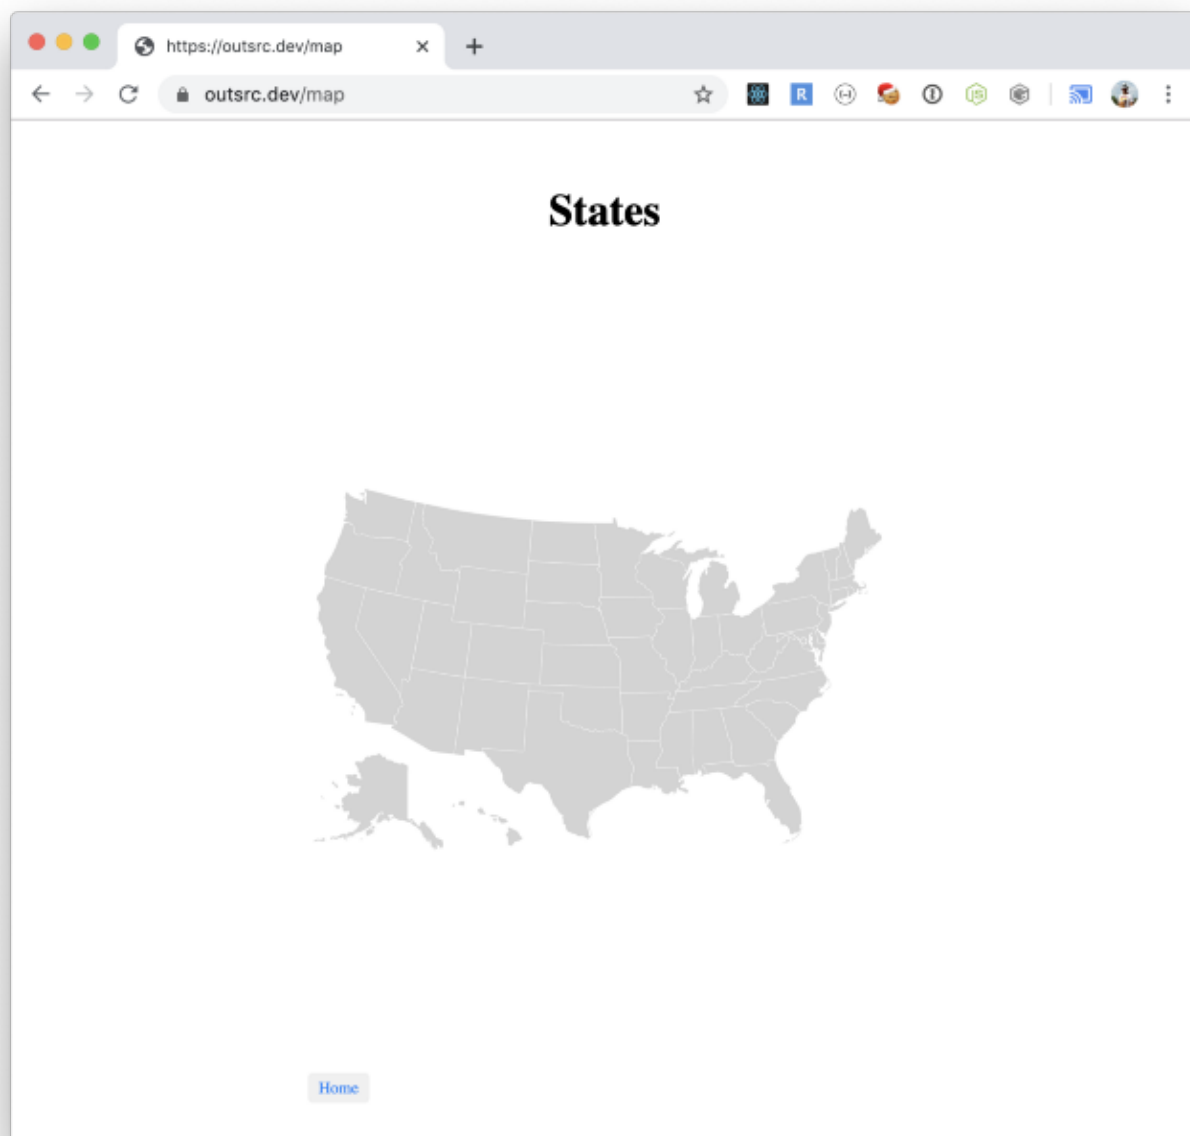
Waiting for deployment "deployment-frontend" rollout to finish: 1
out of 2 new replicas have been updated...
Waiting for deployment "deployment-frontend" rollout to finish: 1
out of 2 new replicas have been updated...
Waiting for deployment "deployment-frontend" rollout to finish: 1
out of 2 new replicas have been updated...
Waiting for deployment "deployment-frontend" rollout to finish: 1
old replicas are pending termination...
Waiting for deployment "deployment-frontend" rollout to finish: 1
old replicas are pending termination...
deployment "deployment-frontend" successfully rolled out
```

Now the deployment's rollout history:

```
$ kubectl rollout history deployment/deployment-frontend
deployment.extensions/deployment-frontend
```

REVISION	CHANGE - CAUSE
1	<none>
2	<none>

And our application has a US Map page now:



New map page deployed

Oh no, we have a bug! rollback..

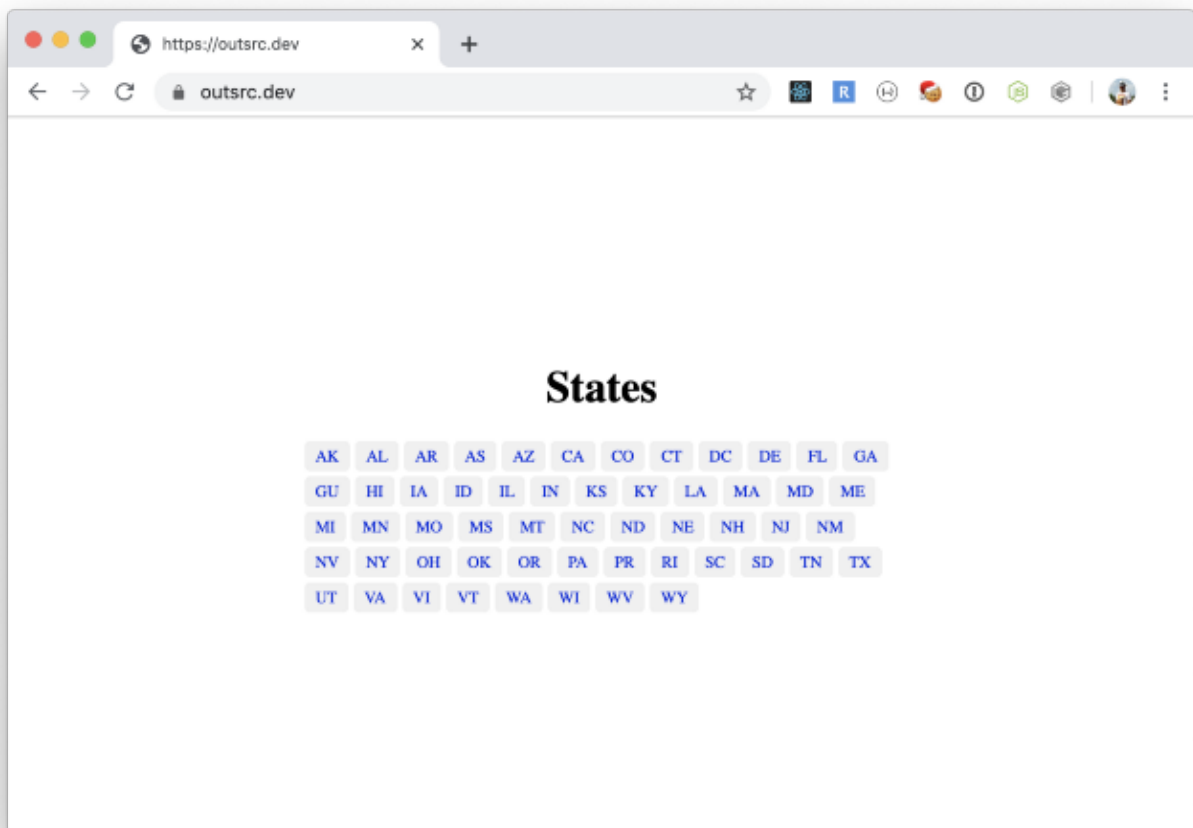
A user found a bug on our newly deployed version of the application. We need to rollback to the previous known working version.

Let's roll it back:


```
$ kubectl rollout undo deployment/deployment-frontend
deployment.extensions/deployment-frontend rolled back
```

```
$ kubectl rollout status -w deployment/deployment-frontend
Waiting for deployment "deployment-frontend" rollout to finish: 1
old replicas are pending termination...
Waiting for deployment "deployment-frontend" rollout to finish: 1
old replicas are pending termination...
deployment "deployment-frontend" successfully rolled out
```

Our application was reverted to previous version.



Application version 1.0.0

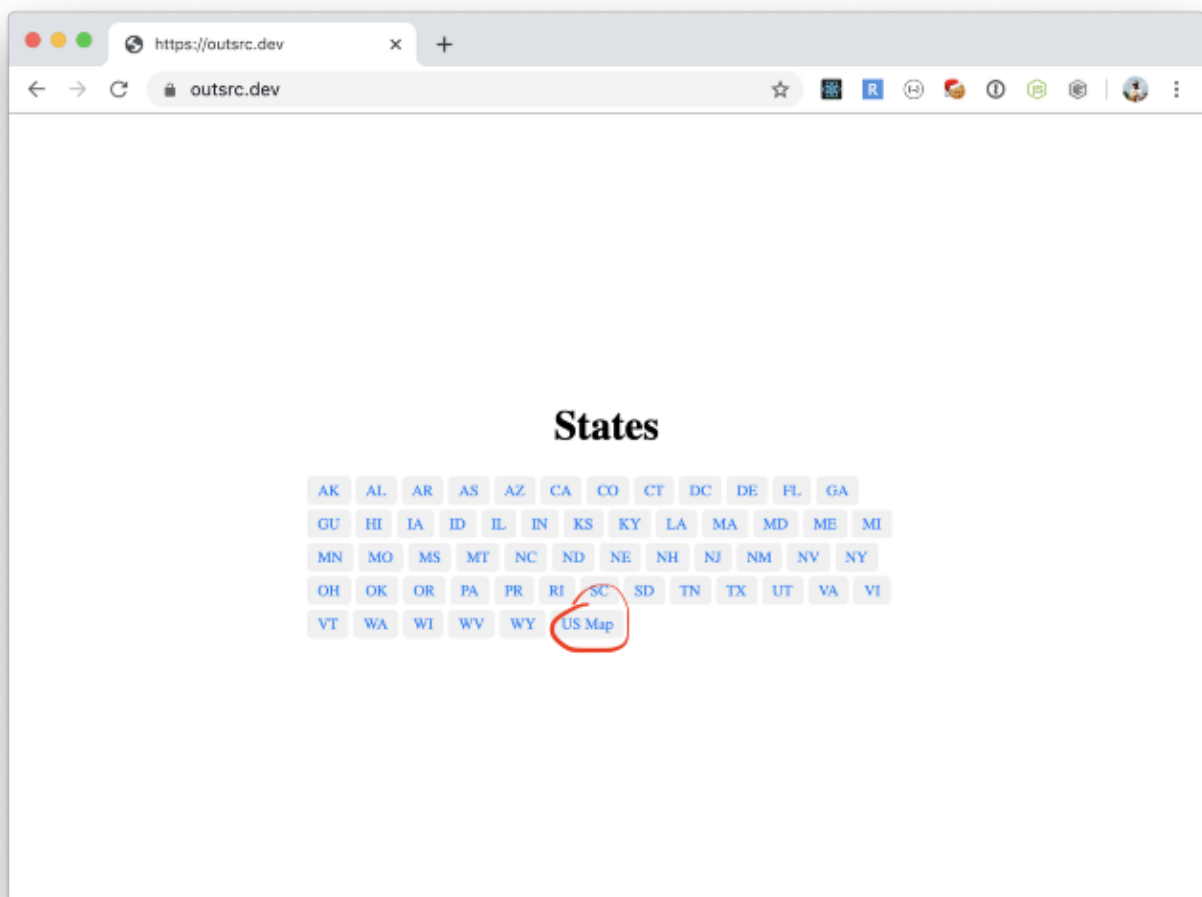
Notice on the history the revision numbers are still going up:

```
$ kubectl rollout history deployment/deployment-frontend
deployment.extensions/deployment-frontend
REVISION  CHANGE-CAUSE
2          <none>
3          <none>
```


We can always revert to any revision. In this case Revision #2 is our Maps revision. Let's bring it back.

```
$ kubectl rollout undo deployment/deployment-frontend --to-revision=2
deployment.extensions/deployment-frontend rolled back
```

```
$ kubectl rollout status -w deployment/deployment-frontend
Waiting for deployment "deployment-frontend" rollout to finish: 1 old replicas are pending termination...
Waiting for deployment "deployment-frontend" rollout to finish: 1 old replicas are pending termination...
deployment "deployment-frontend" successfully rolled out
```



Back to version 1.1.0 (with US Map)

So far so good. Now, this going back and forth on a live website for features and bugs is not a good thing. Our users will feel frustrated if we roll out a feature just to find it has bugs and then roll it back. That's one of the reasons we have different **deployment environments**.

Is very common to find this set of environment:

master | staging | production

- **master:** Or development, usually most updated version, matches the master branch on the repo.
- **staging:** Most close to production, usually where last QA is performed.
- **production:** Is what your users see.

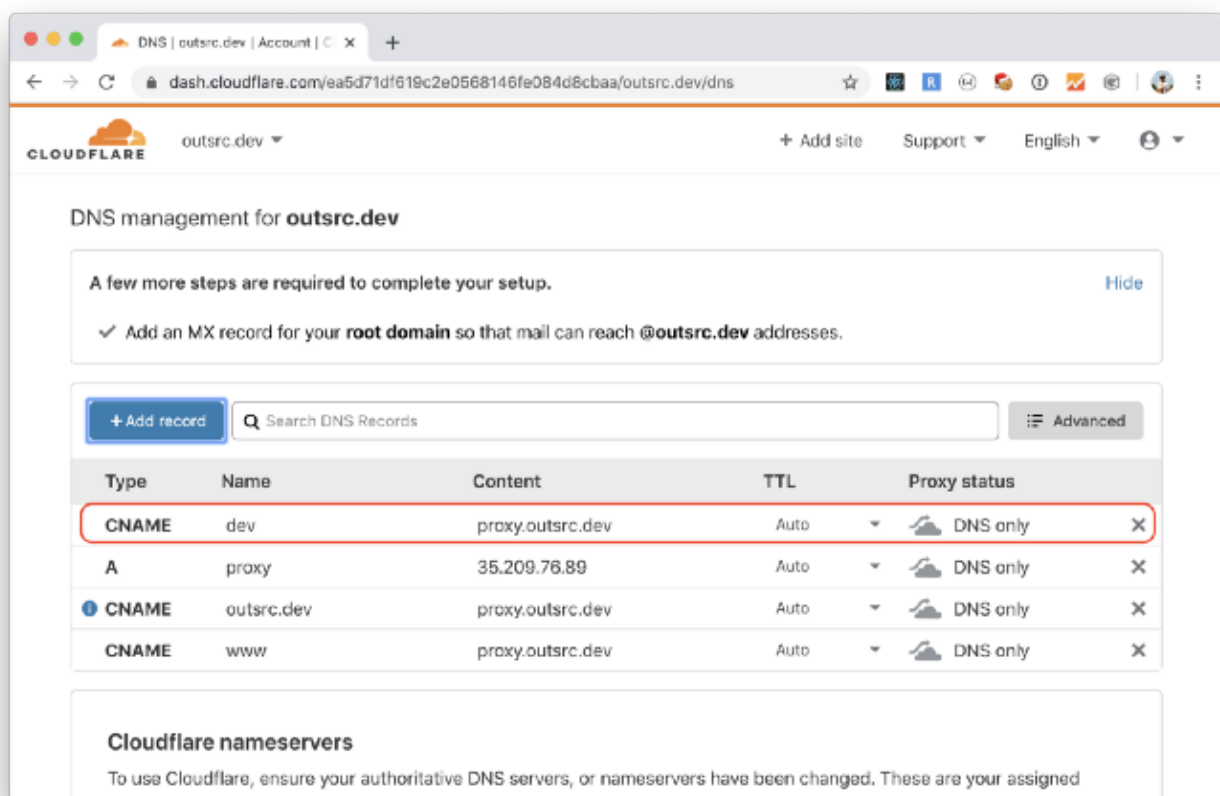
In this setting, any new feature or bugfix will go from master, to staging and then to production.

Out our US States application let's create one more environment: **development**

First we need to select a subdomain, for my current application I will choose:

dev.outsrc.dev (master.outsrc.dev is fine too)

First: DNS, let's make a DNS registry making our subdomain pointing to the cluster proxy IP this is: (Remember our DNS from previous post was hosted on Cloudflare)



dev.outsrc.dev -> CNAME proxy.outsrc.dev

After this a NSLOOKUP of dev.outsrc.dev will return:

```
$ dig dev.outsrc.dev
; <<>> DiG 9.10.6 <<>> dev.outsrc.dev
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 38625
;; flags: qr rd ra; QUERY: 1, ANSWER: 2, AUTHORITY: 0, ADDITIONAL: 1

;; OPT PSEUDOSECTION:
; EDNS: version: 0, flags:; udp: 4096
;; QUESTION SECTION:
;dev.outsrc.dev.      IN A

;; ANSWER SECTION:
dev.outsrc.dev. 300 IN CNAME proxy.outsrc.dev.
proxy.outsrc.dev. 300 IN A 35.209.76.89

;; Query time: 66 msec
;; SERVER: 192.168.1.254#53(192.168.1.254)
;; WHEN: Sun Dec 15 14:48:42 EST 2019
;; MSG SIZE rcvd: 79
```

dev.outsrc.dev points right at our cluster.

Now Lets create a different set of resource files, on a different namespace: outsrc-dev

```
1  apiVersion: v1
2  kind: Namespace
3  metadata:
4    name: outsrc-dev
5  ---
6  apiVersion: extensions/v1beta1
7  kind: Ingress
8  metadata:
9    name: dev-outsrc-dev-ingress
10   namespace: outsrc-dev
11   annotations:
12     kubernetes.io/ingress.class: kong
13     kubernetes.io/tls-acme: 'true'
14     cert-manager.io/cluster-issuer: letsencrypt-production
15   spec:
16     tls:
17       - secretName: dev-outsrc-dev-tls
18     hosts:
```

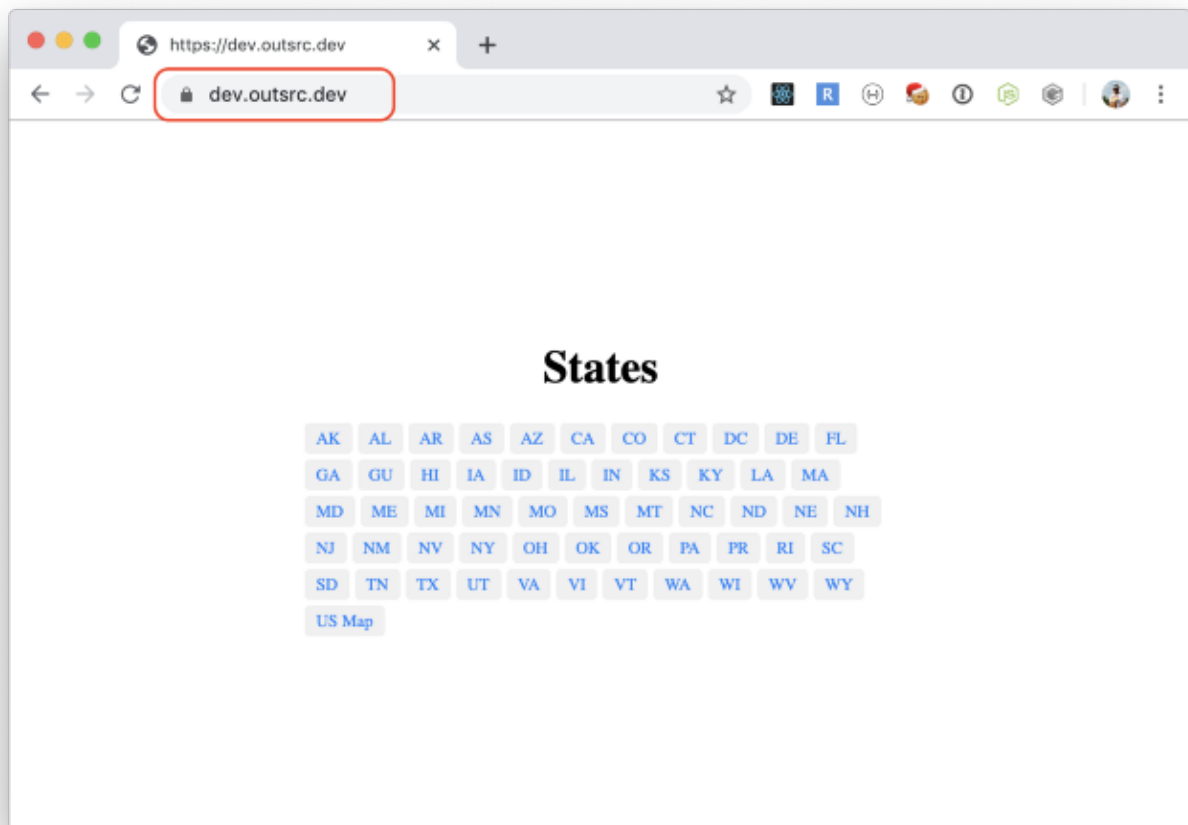
```
19         - dev.outsrc.dev
20     rules:
21     - host: dev.outsrc.dev
22       http:
23         paths:
24         - path: /api
25           backend:
26             serviceName: service-backend
27             servicePort: 3000
28         - path: /
29           backend:
30             serviceName: service-frontend
31             servicePort: 3000
32     ---
33     apiVersion: v1
34     kind: Service
35     metadata:
36       name: service-frontend
37       namespace: outsrc-dev
38       labels:
39         service: front
40     spec:
41       selector:
42         service: front
43       ports:
44       - port: 3000
45         protocol: TCP
46         targetPort: 3000
47     ---
48     apiVersion: apps/v1
49     kind: Deployment
50     metadata:
51       name: deployment-frontend
52       namespace: outsrc-dev
53       labels:
54         service: front
55     spec:
56       replicas: 2
57       selector:
58         matchLabels:
59           service: front
60       template:
61         metadata:
62           labels:
63             service: front
64         spec:
65           containers:
66           - name: frontend-container
```

```
67         image: 'gcr.io/outsrc/outsrc-demo-front:1.1.0'
68         imagePullPolicy: Always
69         ports:
70         - containerPort: 3000
71         env:
72         - name: API_URL
73           value: 'https://outsrc.dev/api'
74         - name: PORT
75           value: '3000'
76 ---
77 apiVersion: v1
78 kind: Service
79 metadata:
80   name: service-backend
81   namespace: outsrc-dev
82   labels:
83     service: back
84 spec:
85   selector:
86     service: back
87   ports:
88   - port: 3000
89     protocol: TCP
90     targetPort: 3000
91 ---
92 apiVersion: apps/v1
93 kind: Deployment
94 metadata:
95   name: deployment-backend
96   namespace: outsrc-dev
97   labels:
98     service: back
99 spec:
100   replicas: 2
101   selector:
102     matchLabels:
103       service: back
104   template:
105     metadata:
106       labels:
107         service: back
108     spec:
109       containers:
110       - name: backend-container
111         image: 'gcr.io/outsrc/outsrc-demo-back:1.0.0'
112         imagePullPolicy: Always
113         ports:
```

```
114         - containerPort: 3000
115     env:
116         - name: PORT
117           value: '3000'
```

```
$ kubectl apply -f outsrc-dev.yml
namespace/outsrc-dev created
ingress.extensions/dev-outsrc-dev-ingress created
service/service-frontend created
deployment.apps/deployment-frontend created
service/service-backend created
deployment.apps/deployment-backend created
```

Ready, we now can access <http://dev.outsrc.dev>



dev.outsrc.dev

Now, this is a subdomain we don't want everybody to be able to access it. This is restricted to the internal developers team, the product managers and QA teams, designers, test engineers, etc.

We need to limit who can access this subdomain. There are several ways we can achieve this. One solution is using Kong Plugins (<https://docs.konghq.com/hub/>), more specific IP Restriction Plugin (<https://docs.konghq.com/hub/kong-inc/ip-restriction/>)

So now you know why I used Kong Ingress on the previous post.

Restrict Access by IP with Kong Plugin

First let's create a plugin resource descriptor

```
1  apiVersion: configuration.konghq.com/v1
2  kind: KongPlugin
3  metadata:
4    name: ip-restriction
5    namespace: outsrc-dev
6  config:
7    whitelist:
8      - 202.110.224.38
9      - 7.107.59.230
10 plugin: ip-restriction
```

ip-restrict.yml hosted with ♥ by GitHub

[view raw](#)

Apply it first:

```
$ kubectl apply -f ip-restrict.yml
kongplugin.configuration.konghq.com/ip-restriction created
```

After this we can modify our Ingress resource file to signal that all routes should be IP restricted:

```
...
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: dev-outsrc-dev-ingress
  namespace: outsrc-dev
  annotations:
    kubernetes.io/ingress.class: kong
    kubernetes.io/tls-acme: 'true'
    cert-manager.io/cluster-issuer: letsencrypt-production
    plugins.konghq.com: ip-restriction
spec:
```

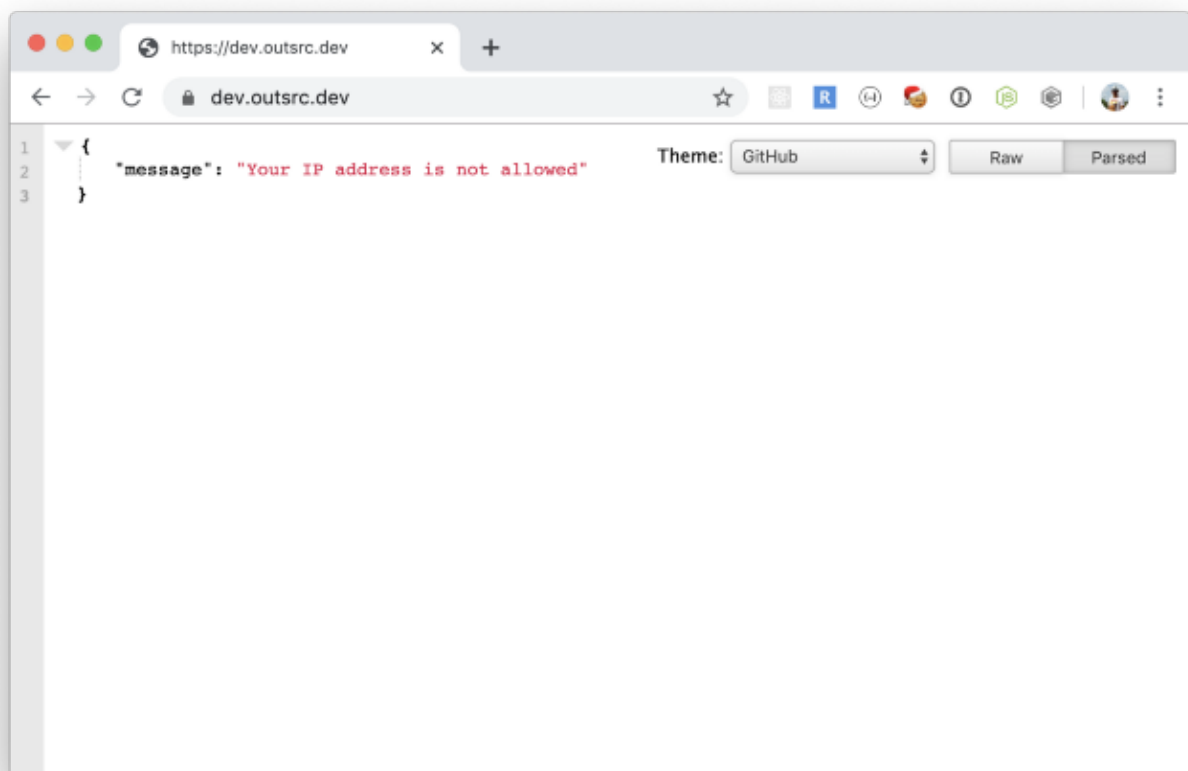


```
tls:
  - secretName: dev-outsrc-dev-tls
    hosts:
      - dev.outsrc.dev
rules:
  - host: dev.outsrc.dev
...
```

And update it:

```
$ kubectl apply -f outsrc-dev.yml
namespace/outsrc-dev unchanged
ingress.extensions/dev-outsrc-dev-ingress configured
service/service-frontend unchanged
deployment.apps/deployment-frontend unchanged
service/service-backend unchanged
deployment.apps/deployment-backend unchanged
```

Now if you try to access from an IP address not whitelisted you will be greeted by:



Denied!

Note: One nice thing about this IP restriction plugin is you can whitelist IP addresses and update only the plugin resource. The Ingress will use the updated list.

Conclusions

- Rolling updates require only a change on the deployment container.
- Rollbacks are actually quite fast.
- We can rollback to any previous deployed version. (Still not sure what are the limits here)
- To avoid going back and forth on production use deployment environments.
- Use different namespaces for your different environments if deploying on the same cluster.
- Optimally, create 2 clusters, separate Production environment from Staging and Development.
- Check Kong Plugins, some of them are really nice or you can write your own (Ready to learn Lua? or send me a text I might help you writing it)

Happy hacking...