

You have 2 free stories left this month. [Sign up and get an extra one for free.](#)

# Traefik cluster as Ingress Controller for Kubernetes



Liejun Tao [Follow](#)

Jun 20, 2019 · 6 min read ★



Traefik cluster as Ingress Controller for Kubernetes


## Introduction

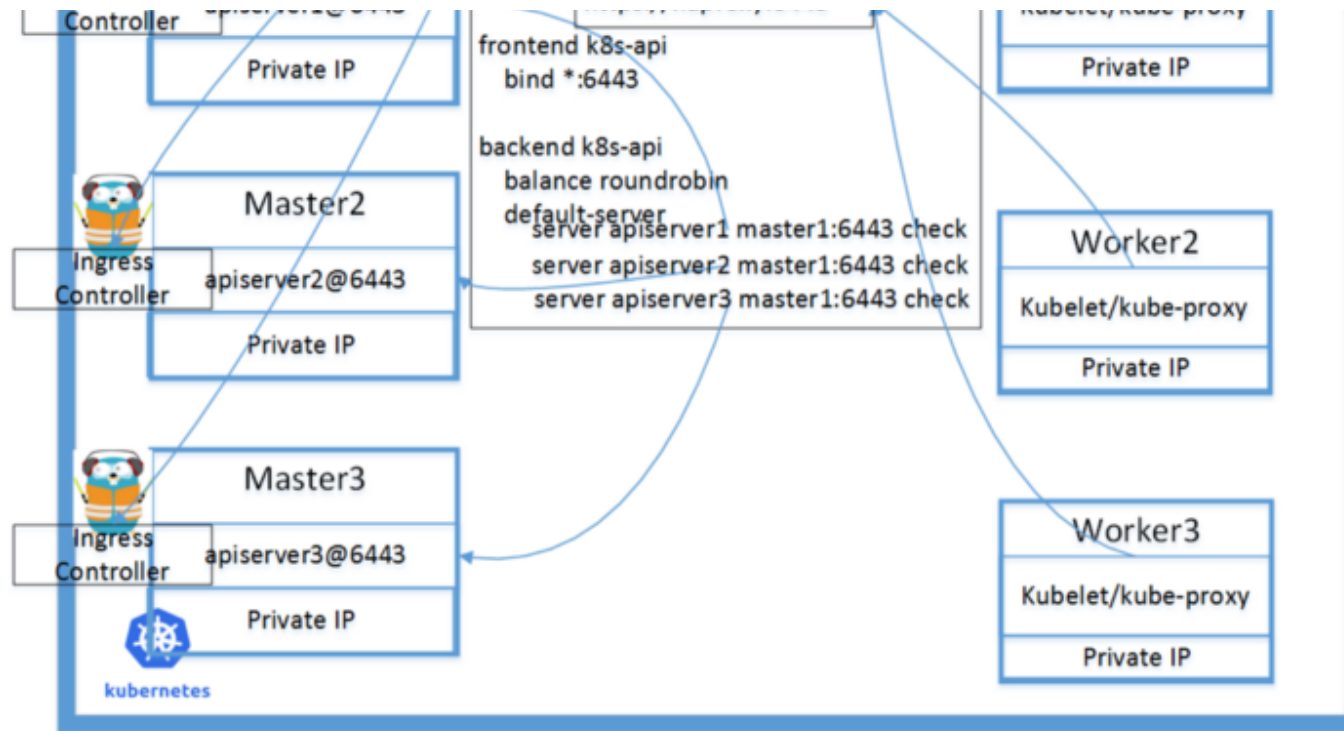
Ingress Controller is the portal to the services running on Kubernetes cluster. To get a highly available cluster, there should be multiple Ingress Controllers working together as a cluster.

Traefik is one of the Ingress Controllers. I use it for its dynamic configuration and automatic LetsEncrypt certificates. There are many instructions to deploy a single Traefik Ingress Controller but not so much details for a Traefik cluster as Ingress Controller. The official document is quite brief, so I'd like to share my experience in this article.

A working HA Kubernetes cluster has at least 3 master nodes and some workers. I have shared my experience [here](#). By adding Traefik cluster, the architecture of the cluster is like below diagram.

Here I deploy Ingress Controllers on the master nodes, which I think is suitable for a low traffic cluster. More instances could be deployed to worker nodes, if the cluster need serve high traffic.





Architecture

Updated 1/9/2020

*Traefik uses Consul Cluster as storage back end. To make it secure/stable, I shared my recent experience about running Consul in Kubernetes for production in [this article](#).*

## Prerequisite

consul command line, example as below

```
cd /tmp && wget  
https://releases.hashicorp.com/consul/1.5.1/consul_1.5.1_linux_amd64.  
zip && unzip consul_1.5.1_linux_amd64.zip && sudo mv consul  
/usr/local/bin
```

traefik command line, example as below

```
cd /tmp && wget  
https://github.com/containous/traefik/releases/download/v1.7.11/trae  
fik_linux-amd64 && chmod +x traefik_linux-amd64 && mv traefik_linux-  
amd64 traefik && sudo mv traefik /usr/local/bin
```

go environment, cfssl, cfssljson

```
# Install gimme  
$ curl -sL -o ~/bin/gimme https://raw.githubusercontent.com/travis-  
ci/gimme/master/gimme  
$ chmod +x ~/bin/gimme
```

```
# Install go
$ eval `gimme stable`

# Install cfssl, cfssljson
go get -u github.com/cloudflare/cfssl/cmd/cfssl
go get -u github.com/cloudflare/cfssl/cmd/cfssljson
```

## Sources used in this article

[github link](#)

## Deploy Consul

Consul is used as KV store for Traefik, while it's actually much more powerful.

I'm following first half of [this tutorial](#) to deploy a 3-replicas consul cluster.

Generate CA and certificate for consul

```
$ cd ca

# Generate CA

$ cfssl gencert -initca config/ca-csr.json | cfssljson -bare ca

# Generate certs

$ cfssl gencert \
  -ca=ca.pem \
  -ca-key=ca-key.pem \
  -config=config/ca-config.json \
  -profile=default \
  config/consul-csr.json | cfssljson -bare consul

$ cfssl gencert \
  -ca=ca.pem \
  -ca-key=ca-key.pem \
  -config=config/ca-config.json \
  -profile=default \
  config/traefik-csr.json | cfssljson -bare traefik
```

## Deploy consul

As Consul need persistent volume, adjust storageClassName in consul/consul\_statefulset.yml properly.

If the cluster has less than 3 worker nodes, remove the “podAntiAffinity” from consul/consul\_statefulset.yml. However to get highly available Consul cluster, it’s better to have 3 or 5 replicas running on different nodes.

```
kubectl apply -f consul/
```

## Generate consul secrets

```
$ cd ca

# Generate gossip key
$ export GOSSIP_ENCRYPTION_KEY=$(consul keygen)

# consul
$ kubectl -n consul create secret generic consul \
  --from-literal="gossip-encryption-key=${GOSSIP_ENCRYPTION_KEY}" \
  --from-file=ca.pem \
  --from-file=consul.pem \
  --from-file=consul-key.pem
```

If things go well, we should have the 3 consul pods up.

```

$ kubectl -n consul get pod
NAME          READY   STATUS    RESTARTS   AGE
consul-0      2/2     Running   0           17s
consul-1      2/2     Running   0           14s
consul-2      2/2     Running   0           11s

# check log
$ kubectl -n consul logs -f consul-0 -c consul
bootstrap_expect > 0: expecting 3 servers
==> Starting Consul agent...
==> Consul agent running!
      Version: 'v1.5.0'
      Node ID: '159adb35-0cfd-2c58-ea51-8e3a6b64ee4c'
      Node name: 'consul-0'
      Datacenter: 'dc1' (Segment: '<all>')
      Server: true (Bootstrap: false)
      Client Addr: [0.0.0.0] (HTTP: 8500, HTTPS: 8443, gRPC: -1,
DNS: 8600)
      Cluster Addr: 10.0.210.154 (LAN: 8301, WAN: 8302)
      Encrypt: Gossip: true, TLS-Outgoing: true, TLS-Incoming:
true

# forward consul to localhost
$ kubectl -n consul port-forward consul-0 8500:8500 &

# run this command in different console to maintain the connection
$ while true; do consul members && sleep 20; done

$ consul members
Node          Address           Status  Type    Build  Protocol  DC
Segment
consul-0      10.0.210.154:8301 alive    server  1.5.0    2          dc1
<all>
consul-1      10.0.143.30:8301  alive    server  1.5.0    2          dc1

```



```
<all>  
consul-2  10.0.181.84:8301  alive   server  1.5.0  2          dc1  
<all>
```

Browse to <http://localhost:8500> for the consul UI



Consul services



Consul nodes



Consul K/V store, now empty

## Import traefik.toml into Consul

I have this traefik.toml.sample as template.

Fill in proper email address and domain name to generate a file traefik.toml.

Import traefik.toml with ‘traefik storeconfig’ command.

```
# Generate from template
$ EMAIL=youremail@example.com DOMAIN=example.com envsubst <
traefik.toml.kv.sample > traefik.toml
```

```
# forward consul to localhost
$ kubectl -n consul port-forward consul-0 8500:8500 &

# magic time
$ traefik storeconfig --consul --consul.endpoint=localhost:8500 --
file.filename=./traefik.toml
...
Writing config to KV
```

If properly imported, check the values in Consul K/V store



K/V store after import

After above importing steps, a manual change for the value of key `traefik/consul/endpoint` is needed. As shown in below screenshot, to 'https://consul.consul.svc.cluster.local:8443', which is the URL to the consul service.



manually modify this valule

If there are existing Let'sEncrypt certificates in a file acme.json, simply uncomment the line

```
#storageFile = "./acme.json"
```

from file traefik.toml and do the import. It will also get imported into Consul.

Then delete the key traefik/acme/account/lock, to allow the Traefik Ingress pods get the lock.

## Deploy Traefik Ingress Controller cluster

Create TLS secret for traefik cluster. It's used to access Consul.

```
$ cd ca
$ kubectl -n kube-system create secret generic traefik-consul \
>     --from-file=ca.pem \
>     --from-file=traefik.pem \
>     --from-file=traefik-key.pem
secret/traefik-consul created
```

Create traefik dashboard secret.

```
kubectl -n kube-system create secret generic kubeseecret --from-file  
auth
```

File 'auth' is created from this command. The username/password will be used to access the traefik dashboard

```
$ htpasswd -c ./auth <username>  
New password:  
Re-type new password:  
Adding password for user testaaa
```

Now deploy it.

The example traefik\_kv.yaml will deploy to master nodes. If this is not desired, adjust the "nodeAffinity" part.

```
DOMAIN=example.com envsubst < traefik_kv.yaml | kubectl apply -f -
```

There are 3 pods for traefik-ingress-controller.

```
$ kubectl -n kube-system get pod -o wide
traefik-ingress-controller-5dlbz          1/1      Running    0
5m21s   10.0.142.69      master3
traefik-ingress-controller-9nxqt          1/1      Running    0
6m34s   10.0.123.196      master2
traefik-ingress-controller-xzkzd          1/1      Running    0
5m34s   10.0.255.14      master1
```

Add labels to the nodes which has Ingress Controller deployed, as I will use the label to handle the cluster reboot sequence. Refer to my article “How to reboot highly available Kubernetes Cluster” ([link](#)).

```
# for all nodes that has Ingress Controller
$ kubectl label nodes master1 node-role.kubernetes.io/ingress-
controller=

$ kubectl get nodes
NAME        STATUS    ROLES                                AGE    VERSION
master1     Ready     ingress-controller,master           49d    v1.14.1
master2     Ready     ingress-controller,master           49d    v1.14.1
master3     Ready     ingress-controller,master           49d    v1.14.1
```

worker1	Ready	<none>	49d	v1.14.1
worker2	Ready	<none>	48d	v1.14.1
worker3	Ready	<none>	48d	v1.14.1

## Add Haproxy rules

When the traefik cluster is ready, add haproxy rules to both primary and backup haproxy hosts. Basically, forward port 80 and 443.

As the DNS for the domain point to the haproxy, now the traefik dashboard <https://traefik.k8s.example.com/> is accessible, with the username/password created above.

## Export acme from consul

I'd like to check/backup the Let'sEncrypt certificates that the Traefik cluster has got for the services. I couldn't find clear instructions and below is my steps after some tries.

```
# Connect to consul
kubectl -n consul port-forward consul-0 8500:8500 &
```



```
# Keep the consul connection open
while true; do consul members && sleep 10; done

# retrieve compressed acme object
consul kv get traefik/acme/account/object > acme.gz

# gunzip
cat acme.gz | gunzip > acme.json

# Format
cat acme.json | jq '.' > acme.json.formatted
```

## Backup/Restore consul

I hit a case that I want to change the Persistent volumes for the consul cluster. So I need destroy the consul cluster and recreate it. Obviously I don't want to lose the existing KV values. I found it is pretty easy with the consul snapshot backup/restore.

1. Create a snapshot as backup
2. Destroy the consul statefulsets and pvc
3. Re-create with new consul configs
4. Restore from snapshot

```
# backup
consul snapshot save backup.snap

# restore
consul snapshot restore backup.snap
```

Thanks for reading.

## Sign up for ITNEXT News - Summer Survey! -

By ITNEXT

3 years ago we started our Medium blog & now we take the next step to unite our community. Therefore, we would like to know more about our readers, authors & their expectations through a survey! [Take a look](#)

 Get this newsletter

Create a free Medium account to get ITNEXT News - Summer Survey! - in your inbox.


Docker

Kubernetes

DevOps

Ingress

Traefik



183 claps





WRITTEN BY

**Liejun Tao**

Follow



**ITNEXT**

ITNEXT is a platform for IT developers & software engineers to share knowledge, connect, collaborate, learn and experience next-gen technologies.

Follow

[See responses \(2\)](#)

## More From Medium


**Creating a Custom Library with Angular**

André Braga in ITNEXT



**Principles, Patterns, and Practices for Effective Infrastructure as Code**

Adarsh Shah in ITNEXT



### Convert paper-based notes to HTML content with Google Vision API

Juan Curti in ITNEXT




### The world's simplest Kubernetes dashboard: k1s

Daniel Weibel in ITNEXT



### Automated HA Kubernetes deployment on Raspberry Pis

Michael Fornaro in ITNEXT



### Should You Switch to Quarkus?

Paul Klinker in ITNEXT



### Go Tutorial: TDD with Go and PostgreSQL [Part II]

Juan Curti in ITNEXT



### Simple way to create kubernetes cluster locally using kind.

Akash Shinde in ITNEXT



## Discover Medium

Welcome to a place where words matter. On Medium, smart voices and original ideas take center stage - with no ads in sight. [Watch](#)

## Make Medium yours

Follow all the topics you care about, and we'll deliver the best stories for you to your homepage and inbox. [Explore](#)

## Become a member

Get unlimited access to the best stories on Medium — and support writers while you're at it. Just \$5/month. [Upgrade](#)

---

# Medium

[About](#)[Help](#)[Legal](#)

